monet-or-not-monet

April 20, 2024

0.1 Monet, or not Monet?

This project investigates the application of Generative Adversarial Networks (GANs) for artistic image creation, specifically targeting the generation of artworks inspired by the Impressionist master, Claude Monet.

GANs are a powerful deep learning architecture comprised of two competing neural networks:

Generator Network: This network acts as the creative force, tasked with generating novel images that adhere to the desired artistic style (Monet in this case). Discriminator Network: Functioning as the art critic, this network aims to accurately distinguish between real Monet paintings and the images produced by the generator network. The training process leverages an adversarial paradigm. The generator network continuously refines its ability to produce art that deceives the discriminator network. Conversely, the discriminator network strives to improve its proficiency in identifying forgeries generated by the generator.

This dynamic relationship fosters a training environment where both networks progressively enhance their performance.

```
[]: import pandas as pd
  import tensorflow as tf
  from tensorflow import keras
  from tensorflow.keras import layers
  #import tensorflow_addons as tfa
  import matplotlib.pyplot as plt
  import numpy as np

'''
import os
  for dirname, _, filenames in os.walk('/kaggle/input'):
      for filename in filenames:
            print(os.path.join(dirname, filename))

'''
```

```
2024-03-08 04:45:15.810971: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered 2024-03-08 04:45:15.811122: E
```

```
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register
    cuFFT factory: Attempting to register factory for plugin cuFFT when one has
    already been registered
    2024-03-08 04:45:16.079659: E
    external/local xla/xla/stream executor/cuda/cuda blas.cc:1515] Unable to
    register cuBLAS factory: Attempting to register factory for plugin cuBLAS when
    one has already been registered
[]: "\nimport os\nfor dirname, _, filenames in os.walk('/kaggle/input'):\n
                                                                               for
     filename in filenames:\n
                                   print(os.path.join(dirname, filename))\n\n"
[]: try:
         tpu = tf.distribute.cluster_resolver.TPUClusterResolver()
        print('Device:', tpu.master())
        tf.config.experimental_connect_to_cluster(tpu)
        tf.tpu.experimental.initialize_tpu_system(tpu)
        strategy = tf.distribute.experimental.TPUStrategy(tpu)
     except:
         strategy = tf.distribute.get_strategy()
     print('Number of replicas:', strategy.num_replicas_in_sync)
     AUTOTUNE = tf.data.experimental.AUTOTUNE
    Number of replicas: 1
[]: MONET_FILENAMES = tf.io.gfile.glob(str('/kaggle/input/gan-getting-started/
     →monet tfrec/*.tfrec'))
     print('Monet TFRecord Files:', len(MONET_FILENAMES))
     PHOTO_FILENAMES = tf.io.gfile.glob(str('/kaggle/input/gan-getting-started/
     →photo tfrec/*.tfrec'))
     print('Photo TFRecord Files:', len(PHOTO_FILENAMES))
    Monet TFRecord Files: 5
    Photo TFRecord Files: 20
[]: IMAGE_SIZE = [256, 256]
     def decode_image(image):
         image = tf.image.decode_jpeg(image, channels=3)
         image = (tf.cast(image, tf.float32) / 127.5) - 1
         image = tf.reshape(image, [*IMAGE_SIZE, 3])
        return image
     def read_tfrecord(example):
        tfrecord_format = {
             "image_name": tf.io.FixedLenFeature([], tf.string),
             "image": tf.io.FixedLenFeature([], tf.string),
```

```
"target": tf.io.FixedLenFeature([], tf.string)
}
example = tf.io.parse_single_example(example, tfrecord_format)
image = decode_image(example['image'])
return image
```

```
[]: def load_dataset(filenames, labeled=True, ordered=False):
    dataset = tf.data.TFRecordDataset(filenames)
    dataset = dataset.map(read_tfrecord, num_parallel_calls=AUTOTUNE)
    return dataset
```

```
[]: monet_ds = load_dataset(MONET_FILENAMES, labeled=True).batch(1) photo_ds = load_dataset(PHOTO_FILENAMES, labeled=True).batch(1)
```

0.1.1 EDA

This is a quick overview of the images and the structure of the data.

The images are 250x250 pixels and can create a 250x250 pixel paintings.

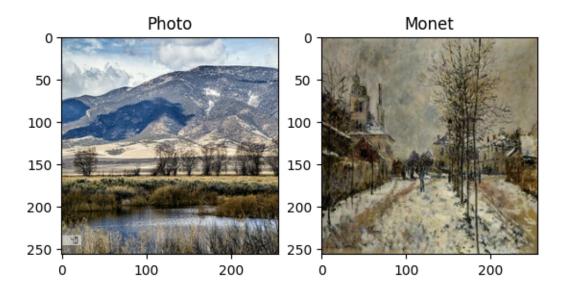
As this is an unsupervised learning algorithm there is less data analysis needed.

```
[]: example_monet = next(iter(monet_ds))
    example_photo = next(iter(photo_ds))

plt.subplot(121)
    plt.title('Photo')
    plt.imshow(example_photo[0] * 0.5 + 0.5)

plt.subplot(122)
    plt.title('Monet')
    plt.imshow(example_monet[0] * 0.5 + 0.5)
```

[]: <matplotlib.image.AxesImage at 0x7e532c0b0d30>



0.1.2 Model Architecture

For this use case a GAN makes the most sense as it takes a generator and a discriminator that train against each other to model a specific style of artwork.

```
[]: def upsample(filters, size, apply_dropout=False):
    initializer = tf.random_normal_initializer(0., 0.02)
    gamma_init = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)

result = keras.Sequential()
    result.add(layers.Conv2DTranspose(filters, size, strides=2,
```

```
[]: def Generator():
         inputs = layers.Input(shape=[256,256,3])
         # bs = batch size
         down stack = [
             downsample(64, 4, apply_instancenorm=False), # (bs, 128, 128, 64)
             downsample(128, 4), # (bs, 64, 64, 128)
             downsample(256, 4), # (bs, 32, 32, 256)
             downsample(512, 4), # (bs, 16, 16, 512)
             downsample(512, 4), # (bs, 8, 8, 512)
             downsample(512, 4), # (bs, 4, 4, 512)
             downsample(512, 4), # (bs, 2, 2, 512)
             downsample(512, 4), # (bs, 1, 1, 512)
         ]
         up_stack = [
             upsample(512, 4, apply_dropout=True), # (bs, 2, 2, 1024)
             upsample(512, 4, apply_dropout=True), # (bs, 4, 4, 1024)
             upsample(512, 4, apply_dropout=True), # (bs, 8, 8, 1024)
             upsample(512, 4), # (bs, 16, 16, 1024)
             upsample(1024,4),
             upsample(512, 4, apply_dropout=True), # (bs, 8, 8, 1024)
             upsample(512, 4), # (bs, 16, 16, 1024)
             upsample(256, 4), # (bs, 32, 32, 512)
             upsample(128, 4), # (bs, 64, 64, 256)
             upsample(64, 4), # (bs, 128, 128, 128)
         ]
         initializer = tf.random_normal_initializer(0., 0.02)
         last = layers.Conv2DTranspose(OUTPUT_CHANNELS, 4,
                                       strides=2,
                                       padding='same',
                                       kernel_initializer=initializer,
                                       activation='tanh') # (bs, 256, 256, 3)
```

```
# Downsampling through the model
skips = []
for down in down_stack:
    x = down(x)
    skips.append(x)

skips = reversed(skips[:-1])

# Upsampling and establishing the skip connections
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = layers.Concatenate()([x, skip])

x = last(x)

return keras.Model(inputs=inputs, outputs=x)
```

```
[]: def Discriminator():
         initializer = tf.random_normal_initializer(0., 0.02)
         gamma_init = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)
         inp = layers.Input(shape=[256, 256, 3], name='input_image')
         x = inp
         down1 = downsample(64, 4, False)(x) # (bs, 128, 128, 64)
         down2 = downsample(128, 4)(down1) # (bs, 64, 64, 128)
         down3 = downsample(256, 4)(down2) # (bs, 32, 32, 256)
         down4 = downsample(512, 4)(down2) # (bs, 32, 32, 256)
         down4 = downsample(1024, 4)(down2) # (bs, 32, 32, 256)
         zero_pad1 = layers.ZeroPadding2D()(down3) # (bs, 34, 34, 256)
         conv = layers.Conv2D(512, 4, strides=1,
                              kernel_initializer=initializer,
                              use_bias=False)(zero_pad1) # (bs, 31, 31, 512)
         norm1 = tf.keras.layers.
      →GroupNormalization(gamma_initializer=gamma_init)(conv)
         leaky_relu = layers.LeakyReLU()(norm1)
         zero_pad2 = layers.ZeroPadding2D()(leaky_relu) # (bs, 33, 33, 512)
         last = layers.Conv2D(1, 4, strides=1,
```

```
kernel_initializer=initializer)(zero_pad2) # (bs, 30,⊔

→30, 1)

return tf.keras.Model(inputs=inp, outputs=last)
```

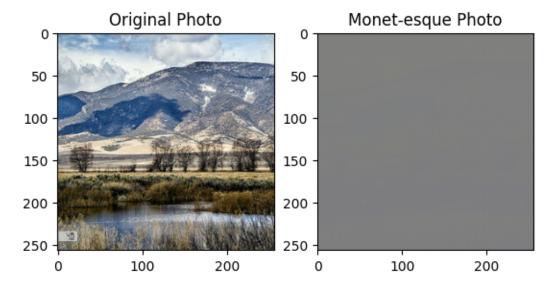
```
[]: with strategy.scope():
    monet_generator = Generator() # transforms photos to Monet-esque paintings
    photo_generator = Generator() # transforms Monet paintings to be more like_
    →photos

monet_discriminator = Discriminator() # differentiates real Monet paintings_
    →and generated Monet paintings
    photo_discriminator = Discriminator() # differentiates real photos and_
    →generated photos
```

```
[]: to_monet = monet_generator(example_photo)

plt.subplot(1, 2, 1)
plt.title("Original Photo")
plt.imshow(example_photo[0] * 0.5 + 0.5)

plt.subplot(1, 2, 2)
plt.title("Monet-esque Photo")
plt.imshow(to_monet[0] * 0.5 + 0.5)
plt.show()
```



```
[]: class CycleGan(keras.Model):
    def __init__(
```

```
self,
    monet_generator,
    photo_generator,
    monet_discriminator,
    photo_discriminator,
    lambda_cycle=10,
):
    super(CycleGan, self).__init__()
    self.m_gen = monet_generator
    self.p_gen = photo_generator
    self.m_disc = monet_discriminator
    self.p_disc = photo_discriminator
    self.lambda_cycle = lambda_cycle
def compile(
    self,
    m_gen_optimizer,
    p_gen_optimizer,
    m_disc_optimizer,
    p_disc_optimizer,
    gen_loss_fn,
    disc_loss_fn,
    cycle_loss_fn,
    identity_loss_fn
):
    super(CycleGan, self).compile()
    self.m_gen_optimizer = m_gen_optimizer
    self.p_gen_optimizer = p_gen_optimizer
    self.m_disc_optimizer = m_disc_optimizer
    self.p_disc_optimizer = p_disc_optimizer
    self.gen_loss_fn = gen_loss_fn
    self.disc_loss_fn = disc_loss_fn
    self.cycle_loss_fn = cycle_loss_fn
    self.identity_loss_fn = identity_loss_fn
def train_step(self, batch_data):
    real_monet, real_photo = batch_data
    with tf.GradientTape(persistent=True) as tape:
        # photo to monet back to photo
        fake_monet = self.m_gen(real_photo, training=True)
        cycled_photo = self.p_gen(fake_monet, training=True)
        # monet to photo back to monet
        fake_photo = self.p_gen(real_monet, training=True)
        cycled_monet = self.m_gen(fake_photo, training=True)
```

```
# generating itself
           same_monet = self.m_gen(real_monet, training=True)
           same_photo = self.p_gen(real_photo, training=True)
           # discriminator used to check, inputing real images
          disc_real_monet = self.m_disc(real_monet, training=True)
          disc_real_photo = self.p_disc(real_photo, training=True)
           # discriminator used to check, inputing fake images
          disc_fake_monet = self.m_disc(fake_monet, training=True)
          disc_fake_photo = self.p_disc(fake_photo, training=True)
           # evaluates generator loss
          monet_gen_loss = self.gen_loss_fn(disc_fake_monet)
          photo_gen_loss = self.gen_loss_fn(disc_fake_photo)
           # evaluates total cycle consistency loss
          total_cycle_loss = self.cycle_loss_fn(real_monet, cycled_monet,_u
self.lambda_cycle) + self.cycle_loss_fn(real_photo, cycled_photo, self.
→lambda_cycle)
           # evaluates total generator loss
          total_monet_gen_loss = monet_gen_loss + total_cycle_loss + self.
dentity_loss_fn(real_monet, same_monet, self.lambda_cycle)
          total_photo_gen_loss = photo_gen_loss + total_cycle_loss + self.
→identity_loss_fn(real_photo, same_photo, self.lambda_cycle)
           # evaluates discriminator loss
          monet_disc_loss = self.disc_loss_fn(disc_real_monet,__

disc_fake_monet)

          photo_disc_loss = self.disc_loss_fn(disc_real_photo,__

¬disc_fake_photo)
       # Calculate the gradients for generator and discriminator
      monet_generator_gradients = tape.gradient(total_monet_gen_loss,
                                                 self.m_gen.

→trainable_variables)
      photo_generator_gradients = tape.gradient(total_photo_gen_loss,
                                                 self.p_gen.
→trainable_variables)
      monet_discriminator_gradients = tape.gradient(monet_disc_loss,
                                                     self.m_disc.
→trainable_variables)
      photo_discriminator_gradients = tape.gradient(photo_disc_loss,
```

```
self.p_disc.
      →trainable_variables)
             # Apply the gradients to the optimizer
             self.m_gen_optimizer.apply_gradients(zip(monet_generator_gradients,
                                                      self.m gen.
      →trainable variables))
             self.p_gen_optimizer.apply_gradients(zip(photo_generator_gradients,
                                                      self.p_gen.
      ⇔trainable_variables))
             self.m_disc_optimizer.apply_gradients(zip(monet_discriminator_gradients,
                                                       self.m_disc.
      ⇔trainable_variables))
             self.p_disc_optimizer.apply_gradients(zip(photo_discriminator_gradients,
                                                       self.p_disc.
      ⇔trainable_variables))
             return {
                 "monet_gen_loss": total_monet_gen_loss,
                 "photo_gen_loss": total_photo_gen_loss,
                 "monet_disc_loss": monet_disc_loss,
                 "photo_disc_loss": photo_disc_loss
             }
[]: with strategy.scope():
         def discriminator_loss(real, generated):
             real_loss = tf.keras.losses.BinaryCrossentropy(from_logits=True,_
      -reduction=tf.keras.losses.Reduction.NONE)(tf.ones_like(real), real)
             generated_loss = tf.keras.losses.BinaryCrossentropy(from_logits=True,_
      -reduction=tf.keras.losses.Reduction.NONE)(tf.zeros_like(generated), ا
      ⇒generated)
             total_disc_loss = real_loss + generated_loss
             return total_disc_loss * 0.5
[]: with strategy.scope():
         def generator_loss(generated):
             return tf.keras.losses.BinaryCrossentropy(from_logits=True,_
      oreduction=tf.keras.losses.Reduction.NONE)(tf.ones_like(generated), generated)
```

```
[]: with strategy.scope():
         def calc_cycle_loss(real_image, cycled_image, LAMBDA):
             loss1 = tf.reduce_mean(tf.abs(real_image - cycled_image))
             return LAMBDA * loss1
[]: with strategy.scope():
         def identity_loss(real_image, same_image, LAMBDA):
             loss = tf.reduce_mean(tf.abs(real_image - same_image))
             return LAMBDA * 0.5 * loss
[]: with strategy.scope():
         monet_generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
         photo_generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
         monet_discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
         photo_discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
[]: with strategy.scope():
         cycle_gan_model = CycleGan(
             monet_generator, photo_generator, monet_discriminator,_
      →photo_discriminator
         cycle_gan_model.compile(
             m_gen_optimizer = monet_generator_optimizer,
            p_gen_optimizer = photo_generator_optimizer,
            m_disc_optimizer = monet_discriminator_optimizer,
             p_disc_optimizer = photo_discriminator_optimizer,
             gen_loss_fn = generator_loss,
             disc_loss_fn = discriminator_loss,
             cycle_loss_fn = calc_cycle_loss,
             identity_loss_fn = identity_loss
         )
[]: cycle_gan_model.fit(
         tf.data.Dataset.zip((monet_ds, photo_ds)),
         epochs=10
    Epoch 1/10
    WARNING: All log messages before absl::InitializeLog() is called are written to
    STDERR
    I0000 00:00:1709873302.333878
                                      128 device_compiler.h:186] Compiled cluster
    using XLA! This line is logged at most once for the lifetime of the process.
    300/300
                        333s 572ms/step -
```

```
monet_disc_loss: 0.5746 - monet_gen_loss: 5.7664 - photo_disc_loss: 0.5516 -
photo_gen_loss: 5.8791 - loss: 0.0000e+00
Epoch 2/10
/opt/conda/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of
data; interrupting training. Make sure that your dataset or generator can
generate at least `steps_per_epoch * epochs` batches. You may need to use the
`.repeat()` function when building your dataset.
 self.gen.throw(typ, value, traceback)
300/300
                   171s 569ms/step -
monet_disc_loss: 0.5941 - monet_gen_loss: 3.9958 - photo_disc_loss: 0.6298 -
photo_gen_loss: 3.9536 - loss: 0.0000e+00
Epoch 3/10
149/300
                   1:26 570ms/step -
monet_disc_loss: 0.6219 - monet_gen_loss: 3.8039 - photo_disc_loss: 0.6301 -
photo_gen_loss: 3.7763
```

0.1.3 Results and Analysis

Results will be created off of an image set that neither the generator or discriniator have seen then ranked by a an algorithm called MiFID. MiFID uses the minimum cosine distance of all training samples in the feature space then averages across all samples to create a score of how well the GAN works to create monet like paintings. The lower the score the better the model. This model has been evaluated at 76.12. The generator and discrimantor both worked well allowing for above average for the kagle competition. I find this to be very encouraging. I think that using more layers and using a deticated GPU would create even better results

0.1.4 Conclusion

Generative Adversial Networks are a very interesting way of creating generative AI artwork. When you have an expressed purpose realatively small use case for a more expansive generative AI solution you would need much more data and layers to the model that could create artwork of any artist or use case. This was a fun and educational way to interact with generative AI as it is becomming a more and more influential part of our daily lives.

Submission

```
[]: _, ax = plt.subplots(5, 2, figsize=(12, 12))
for i, img in enumerate(photo_ds.take(5)):
    prediction = monet_generator(img, training=False)[0].numpy()
    prediction = (prediction * 127.5 + 127.5).astype(np.uint8)
    img = (img[0] * 127.5 + 127.5).numpy().astype(np.uint8)

ax[i, 0].imshow(img)
    ax[i, 1].imshow(prediction)
    ax[i, 0].set_title("Input Photo")
    ax[i, 1].set_title("Monet-esque")
    ax[i, 0].axis("off")
    ax[i, 1].axis("off")
```

```
plt.show()

[]: import PIL
   ! mkdir ../images

[]: import shutil
   shutil.make_archive("/kaggle/working/images", 'zip', "/kaggle/images")

[]: i = 1
   for img in photo_ds:
        prediction = monet_generator(img, training=False)[0].numpy()
        prediction = (prediction * 127.5 + 127.5).astype(np.uint8)
        im = PIL.Image.fromarray(prediction)
        im.save("../images/" + str(i) + ".jpg")
        i += 1
```